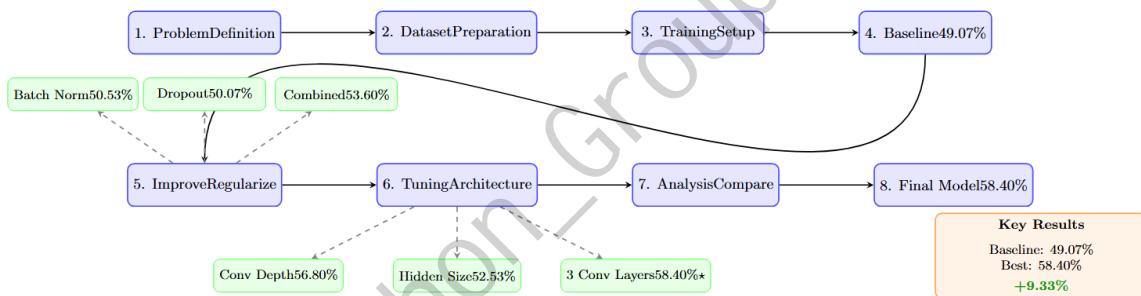


# ARIN5101 Course Project: CNN-based Image Classification for Fine-Grained Mammal Recognition

## Group Information

| Item         | Details                            |
|--------------|------------------------------------|
| Group Number | 20                                 |
| Team Members | Huang Junqin, Rao Ridi, Li Shuying |
| Course       | ARIN5101                           |
| Date         | November 27th, 2025                |

## CNN Image Classification Project Flow



## **1. Problem Definition and Objective**

## **1.1 Problem Statement**

This project addresses the challenge of **fine-grained image classification** - distinguishing between visually similar animal species using Convolutional Neural Networks (CNNs). Fine-grained classification is particularly challenging because categories share many visual characteristics while having subtle distinguishing features.

## 1.2 Objective

Build and systematically compare CNN models for mammal classification, exploring:

- Baseline CNN architecture
  - Regularization techniques (Dropout, Batch Normalization)
  - Architecture tuning (layer depth, width, number of layers)

## 2. Dataset Description

## 2.1 Data Source

We use a subset of the **CIFAR-100 dataset**, focusing on 15 mammal classes organized into three superclasses:

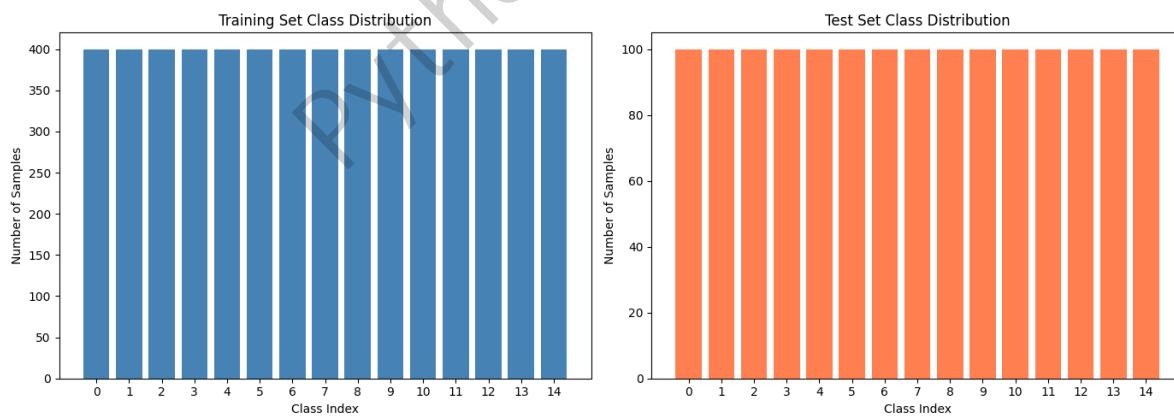
| Superclass              | Classes                                 |
|-------------------------|---|
| <b>Large Carnivores</b> | bear, leopard, lion, tiger, wolf        |
| <b>Medium Mammals</b>   | fox, porcupine, opossum, raccoon, skunk |
| <b>Small Mammals</b>    | hamster, mouse, rabbit, shrew, squirrel |

## 2.2 Dataset Statistics

| Metric           | Value                 |
|------------------|-----------------------|
| Total Classes    | 15                    |
| Training Images  | 7,500 (500 per class) |
| Test Images      | 1,500 (100 per class) |
| Image Resolution | 32 × 32 × 3 (RGB)     |
| Color Channels   | 3 (Red, Green, Blue)  |

## 2.3 Data Distribution

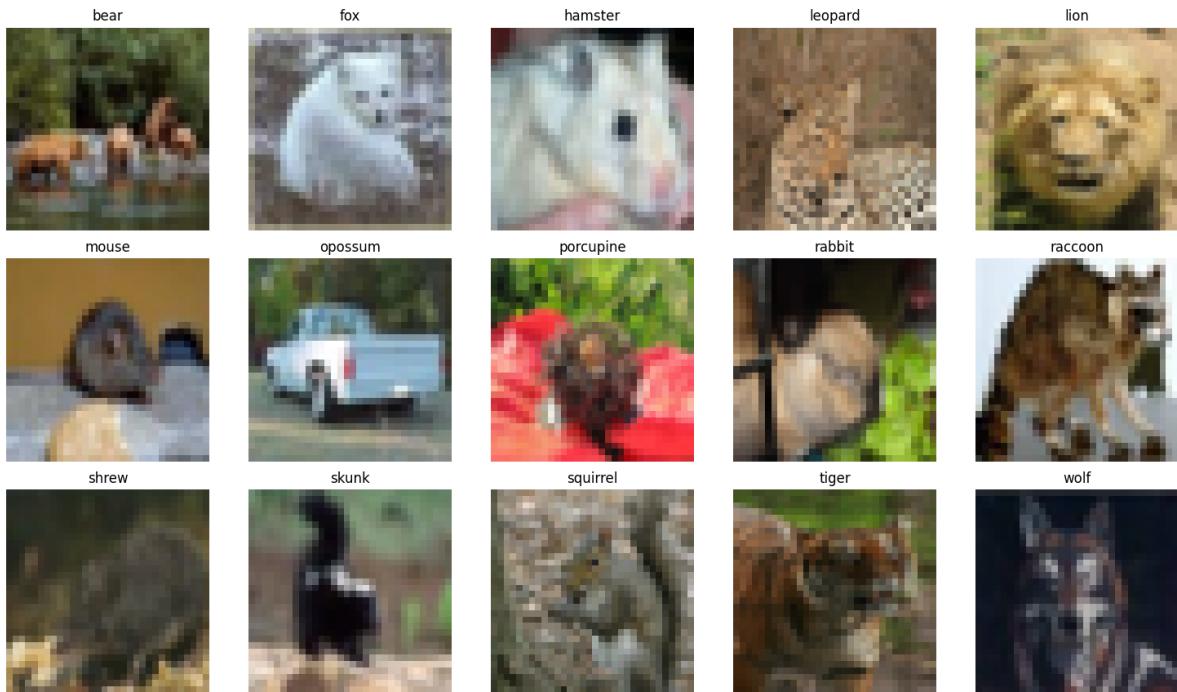
The dataset is **balanced** across all classes, with approximately equal samples per category. This eliminates class imbalance as a confounding factor in model evaluation.



## 2.4 Sample Images

Sample images from each class demonstrate the challenge of fine-grained classification - many mammals share similar body shapes, fur textures, and color patterns.

Sample Images from Each Class



## 2.5 Data Preprocessing

```
# Normalization: Scale pixel values to [0, 1]
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# Convert to PyTorch format (NCHW)
X_tensor = torch.FloatTensor(X).permute(0, 3, 1, 2)
```

### Preprocessing Steps:

1. Load CIFAR-100 dataset using `torchvision.datasets`
2. Filter to 15 mammal classes
3. Remap labels to 0-14
4. Normalize pixel values to [0, 1] range
5. Convert to PyTorch tensors (NCHW format)
6. Create DataLoader with batch size 64

## 3. Training Setup and Validation Strategy

### 3.1 Dataset Splitting Strategy

| Split      | Samples | Percentage |
|------------|---------|------------|
| Training   | 6,000   | 80%        |
| Validation | 1,500   | 20%        |
| Test       | 1,500   | Held out   |

**Stratification:** We use stratified splitting to ensure each class is proportionally represented in both training and validation sets.

```

x_train_split, x_val, y_train_split, y_val = train_test_split(
    X_train, y_train,
    test_size=0.2,
    random_state=0,
    stratify=y_train # Ensures balanced class distribution
)

```

## 3.2 Training Configuration

| Parameter               | Value            |
|-------------------------|------------------|
| Optimizer               | Adam             |
| Loss Function           | CrossEntropyLoss |
| Batch Size              | 64               |
| Max Epochs              | 50               |
| Early Stopping Patience | 10 epochs        |
| Learning Rate           | Default (0.001)  |

## 3.3 Early Stopping Implementation

To prevent overfitting, we implement early stopping that:

- Monitors validation loss
- Saves best model weights
- Stops training if no improvement for 10 consecutive epochs
- Restores best weights after stopping

```

# Early stopping logic
if val_loss < best_val_loss:
    best_val_loss, wait = val_loss, 0
    best_weights = model.state_dict().copy()
else:
    wait += 1
    if wait >= patience:
        model.load_state_dict(best_weights)
        break

```

## 3.4 Metrics Tracked

- **Training Loss** - Cross-entropy loss on training batches
- **Training Accuracy** - Classification accuracy on training set
- **Validation Loss** - Cross-entropy loss on validation set
- **Validation Accuracy** - Classification accuracy on validation set
- **Test Accuracy** - Final evaluation metric

## 4. Baseline Model Architecture

---

## 4.1 Architecture Design

```
INPUT (32x32x3)

Conv Layer 1: 32 filters, 3x3 kernel, padding=1
Activation: ReLU
MaxPooling: 2x2
Output: 16x16x32

Conv Layer 2: 64 filters, 3x3 kernel, padding=1
Activation: ReLU
MaxPooling: 2x2
Output: 8x8x64

Flatten: 8x8x64 = 4,096 features

Dense Layer: 128 units, ReLU

Output Layer: 15 units (one per class)
```

## 4.2 PyTorch Implementation

```
class BaselineCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(32, 64, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2),
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64 * 8 * 8, 128), nn.ReLU(),
            nn.Linear(128, 15)
        )

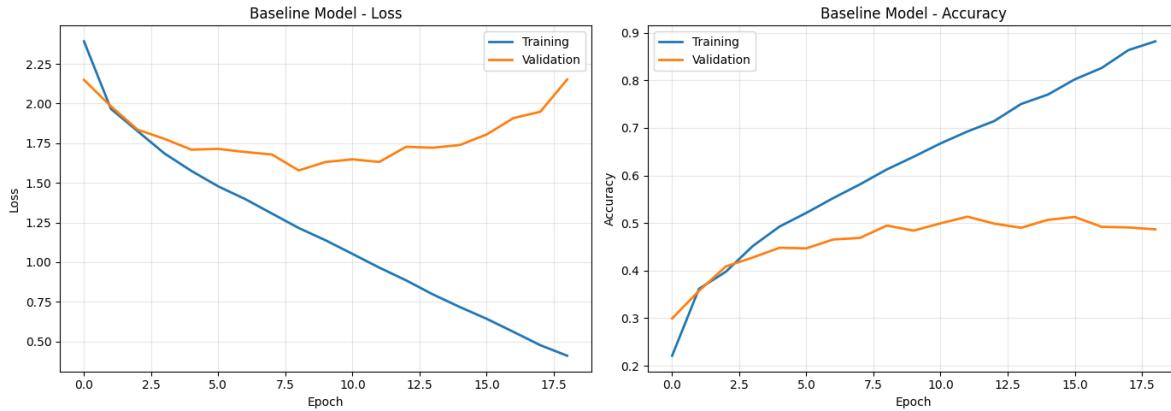
    def forward(self, x):
        return self.classifier(self.features(x))
```

## 4.3 Model Parameters

| Component    | Parameters                                     |
|--------------|--|
| Conv Layer 1 | $3 \times 3 \times 3 \times 32 + 32 = 896$     |
| Conv Layer 2 | $3 \times 3 \times 32 \times 64 + 64 = 18,496$ |
| Dense Layer  | $4,096 \times 128 + 128 = 524,416$             |
| Output Layer | $128 \times 15 + 15 = 1,935$                   |
| <b>Total</b> | <b>545,743</b>                                 |

## 4.4 Baseline Performance

| Metric        | Value         |
|---------------|---------------|
| Test Accuracy | <b>49.07%</b> |
| Test Loss     | 1.5958        |



### Observations:

- Training accuracy reaches ~70% while validation plateaus at ~50%
- Clear signs of overfitting (gap between training and validation)
- Model learns basic features but struggles with fine-grained distinctions

## 5. Model Improvements

We implement two regularization techniques:

1. **Dropout** - Prevents overfitting by randomly dropping connections
2. **Batch Normalization** - Stabilizes training by normalizing layer inputs

### 5.1 Dropout Implementation

#### Dropout Rates:

- After each conv layer: 25%
- Before output layer: 50%

```
class DropoutCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1), nn.ReLU(),
            nn.MaxPool2d(2), nn.Dropout2d(0.25), # 25% dropout
            nn.Conv2d(32, 64, 3, padding=1), nn.ReLU(),
            nn.MaxPool2d(2), nn.Dropout2d(0.25),
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64*8*8, 128), nn.ReLU(),
            nn.Dropout(0.5), # 50% dropout
            nn.Linear(128, 15)
        )
```

### Dropout Results:

| Metric        | Value         |
|---------------|---------------|
| Test Accuracy | <b>50.07%</b> |
| Improvement   | +1.00%        |

## 5.2 Batch Normalization Implementation

```
class BatchNormCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1),
            nn.BatchNorm2d(32), # Normalize after conv
            nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(32, 64, 3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(), nn.MaxPool2d(2),
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64*8*8, 128),
            nn.BatchNorm1d(128), # Normalize dense layer
            nn.ReLU(),
            nn.Linear(128, 15)
    )
```

### Batch Normalization Results:

| Metric        | Value         |
|---------------|---------------|
| Test Accuracy | <b>50.53%</b> |
| Improvement   | +1.46%        |

## 5.3 Combined Improvements (Dropout + Batch Normalization)

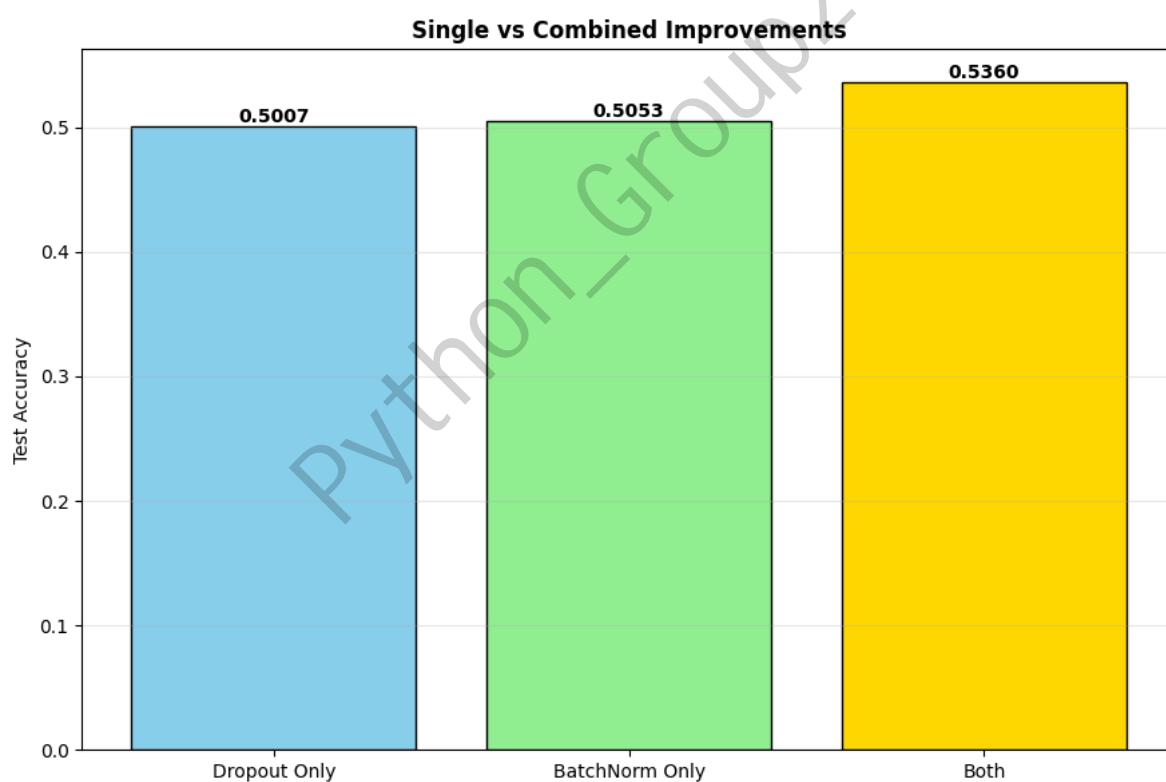
```
class BothCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1), nn.BatchNorm2d(32),
            nn.ReLU(), nn.MaxPool2d(2), nn.Dropout2d(0.25),
            nn.Conv2d(32, 64, 3, padding=1), nn.BatchNorm2d(64),
            nn.ReLU(), nn.MaxPool2d(2), nn.Dropout2d(0.25),
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64*8*8, 128), nn.BatchNorm1d(128),
            nn.ReLU(), nn.Dropout(0.5),
            nn.Linear(128, 15)
```

## Combined Results:

| Metric        | Value  |
|---------------|--------|
| Test Accuracy | 53.60% |
| Improvement   | +4.53% |

## 5.4 Improvement Techniques Comparison

| Model                | Test Accuracy | Improvement   |
|----------------------|---------------|---------------|
| Baseline             | 49.07%        | -             |
| Dropout Only         | 50.07%        | +1.00%        |
| BatchNorm Only       | 50.53%        | +1.46%        |
| <b>Both Combined</b> | <b>53.60%</b> | <b>+4.53%</b> |



**Key Finding:** Combining both techniques provides synergistic benefits - Dropout prevents co-adaptation of features while BatchNorm stabilizes the gradients during training.

## 6. Model Tuning Experiments

### 6.1 Convolution Layer Depth (64, 128)

Increasing filter counts from (32, 64) to (64, 128):

```
depth_model = TunedCNN(conv_depths=(64, 128))
```

| Metric        | Value         |
|---------------|---------------|
| Test Accuracy | <b>56.80%</b> |
| Improvement   | +7.73%        |

## 6.2 Hidden Layer Size (256 units)

Increasing hidden layer from 128 to 256 units:

```
hidden_model = TunedCNN(hidden_units=256)
```

| Metric        | Value         |
|---------------|---------------|
| Test Accuracy | <b>52.53%</b> |
| Improvement   | +3.46%        |

## 6.3 Three Convolutional Layers

Adding a third conv layer with 128 filters:

```
class TunedCNN(nn.Module):
    def __init__(self, conv_depths=(32, 64), hidden_units=128,
three_conv=False):
        super().__init__()
        layers_list = [
            nn.Conv2d(3, conv_depths[0], 3, padding=1),
            nn.BatchNorm2d(conv_depths[0]),
            nn.ReLU(), nn.MaxPool2d(2), nn.Dropout2d(0.25),
            nn.Conv2d(conv_depths[0], conv_depths[1], 3, padding=1),
            nn.BatchNorm2d(conv_depths[1]),
            nn.ReLU(), nn.MaxPool2d(2), nn.Dropout2d(0.25),
        ]

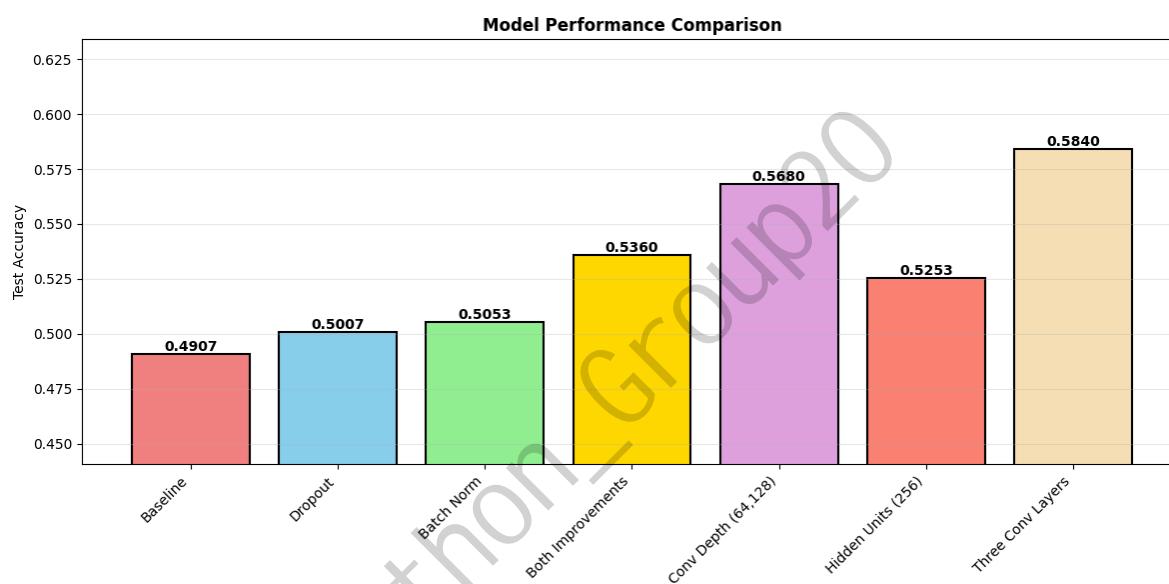
        if three_conv:
            layers_list += [
                nn.Conv2d(conv_depths[1], 128, 3, padding=1),
                nn.BatchNorm2d(128),
                nn.ReLU(), nn.MaxPool2d(2), nn.Dropout2d(0.25)
            ]
```

| Metric        | Value         |
|---------------|---------------|
| Test Accuracy | <b>58.40%</b> |
| Improvement   | +9.33%        |

## 7. Performance Analysis

## 7.1 Complete Results Summary

| Model                    | Test Accuracy | Test Loss     | Rank     |
|--------------------------|---------------|---------------|----------|
| Baseline                 | 49.07%        | 1.5958        | 7        |
| Dropout                  | 50.07%        | 1.5691        | 6        |
| Batch Norm               | 50.53%        | 1.5829        | 5        |
| Both Improvements        | 53.60%        | 1.4348        | 4        |
| Hidden Units (256)       | 52.53%        | 1.4700        | 5        |
| Conv Depth (64,128)      | 56.80%        | 1.3754        | 2        |
| <b>Three Conv Layers</b> | <b>58.40%</b> | <b>1.2970</b> | <b>1</b> |



## 7.2 Training Curves Analysis

### Improvements Comparison:

- Baseline shows clear overfitting with diverging train/val curves
- Dropout reduces overfitting but slightly slows convergence
- BatchNorm accelerates early training but still overfits
- Combined approach shows best balance

### Tuning Comparison:

- Deeper conv layers (64, 128) improve feature extraction
- Larger hidden layer (256) shows diminishing returns
- Three conv layers achieve best generalization

## 7.3 Loss Analysis

| Model             | Test Loss | Change from Baseline |
|-------------------|-----------|----------------------|
| Baseline          | 1.5958    | -                    |
| Three Conv Layers | 1.2970    | -18.7%               |

Lower loss indicates more confident and accurate predictions.

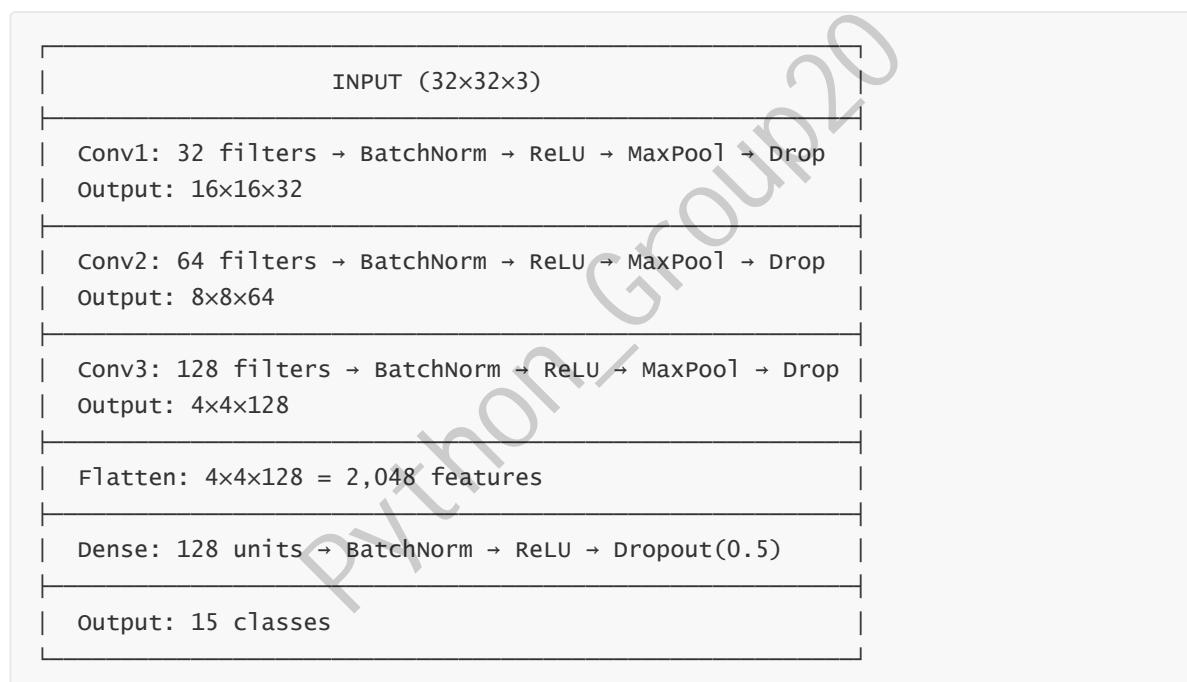
## 8. Final Model Summary and Reflections

### 8.1 Best Performing Model

Three Convolutional Layers with Dropout and Batch Normalization

| Metric                    | Value            |
|---------------------------|------------------|
| Architecture              | 3 Conv + 1 Dense |
| Test Accuracy             | <b>58.40%</b>    |
| Improvement over Baseline | <b>+9.33%</b>    |
| Test Loss                 | 1.2970           |

### 8.2 Architecture Summary



### 8.3 Why Three Conv Layers Performs Best

#### 1. Hierarchical Feature Learning:

- Layer 1: Low-level features (edges, textures)
- Layer 2: Mid-level features (patterns, shapes)
- Layer 3: High-level features (body parts, distinctive features)

#### 2. Progressive Abstraction:

- Each pooling layer reduces spatial dimensions
- Deeper layers capture more abstract, class-specific information

#### 3. Better Regularization:

- More opportunities for dropout
- BatchNorm at each layer stabilizes training

#### 4. Appropriate Complexity:

- Sufficient capacity to learn fine-grained distinctions

- Not too deep to cause vanishing gradients on small images

## 8.4 Key Findings

| Finding                             | Evidence   |
|-------------------------------------|--|
| Regularization is crucial           | Combined Dropout+BatchNorm outperforms single techniques |
| Depth matters more than width       | 3 conv layers > larger filters or hidden units           |
| Early stopping prevents overfitting | Best models stopped before 50 epochs                     |
| Feature hierarchy is important      | More layers = better abstract feature learning           |

## 8.5 Lessons Learned

1. **Start Simple:** Baseline model helps identify overfitting issues
2. **Combine Techniques:** Synergistic effects from Dropout + BatchNorm
3. **Depth Over Width:** Adding layers more effective than wider layers
4. **Monitor Validation:** Early stopping crucial for generalization
5. **Systematic Comparison:** Testing one variable at a time isolates effects

## 8.6 Limitations and Future Work

### Limitations:

- 58.40% accuracy still leaves room for improvement
- Small image size (32×32) limits feature detail
- Simple augmentation not explored

### Future Improvements:

- Data augmentation (rotation, flipping, color jittering)
- Transfer learning with pre-trained models (ResNet, VGG)
- Learning rate scheduling
- More sophisticated architectures (skip connections, attention)

## 9. Confusion Matrix & Classification Report

To gain deeper insights into model performance across individual classes, we analyze the confusion matrix and per-class metrics of our best performing model (Three Conv Layers).

### 9.1 Implementation

```
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns

# Evaluate best model on test set
best_model.eval()
all_preds, all_labels = [], []

with torch.no_grad():
    for x, y in test_loader:
```

```

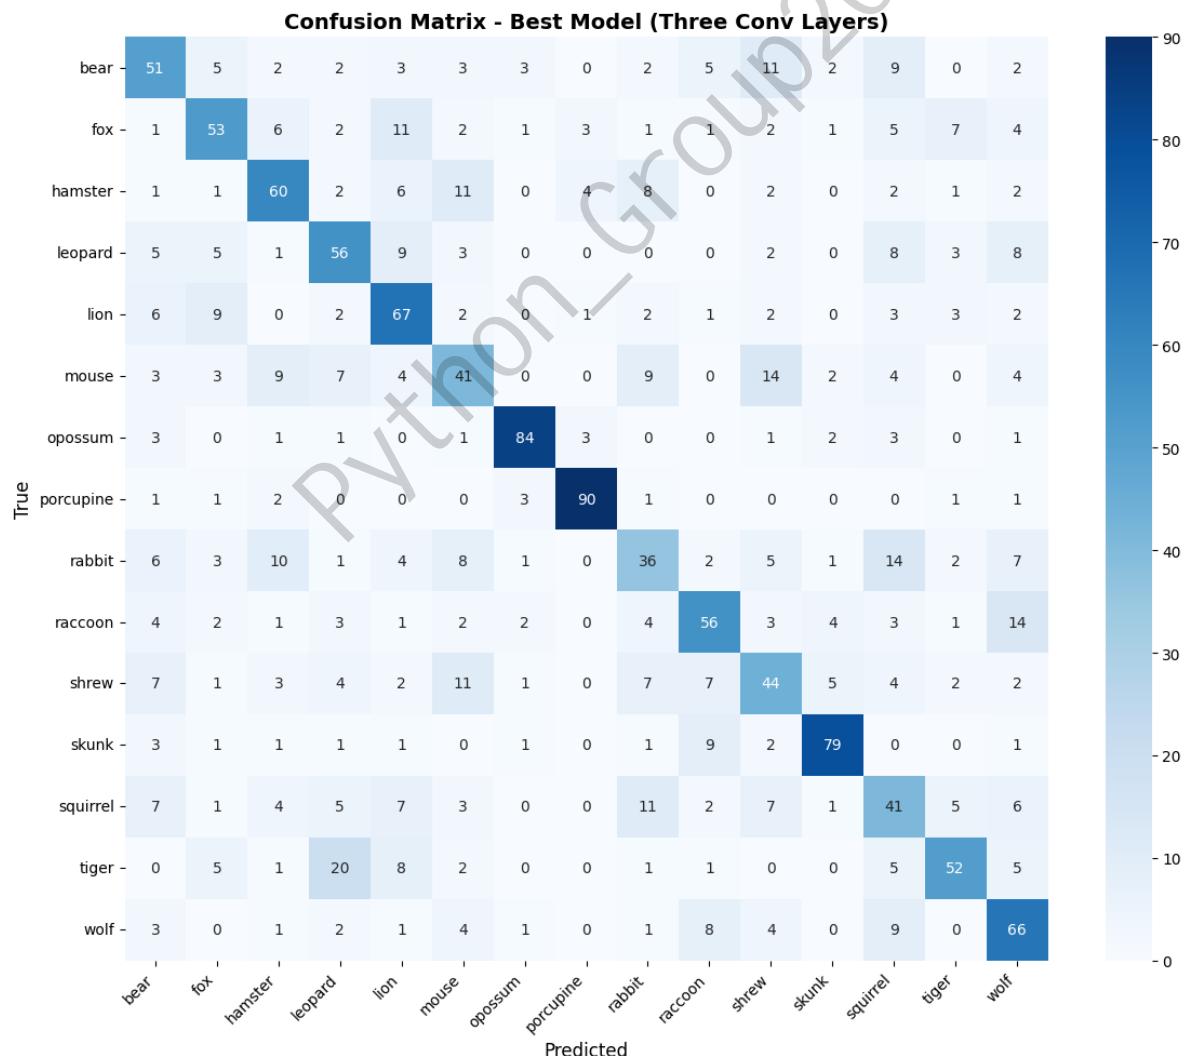
X = X.to(device)
outputs = best_model(X)
preds = outputs.argmax(dim=1).cpu().numpy()
all_preds.extend(preds)
all_labels.extend(y.numpy())

# Generate classification report
print(classification_report(all_labels, all_preds, target_names=class_names))

# Plot confusion matrix
cm = confusion_matrix(all_labels, all_preds)
plt.figure(figsize=(12, 10))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.title('Confusion Matrix - Best Model (Three Conv Layers)')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.tight_layout()
plt.show()

```

## 9.2 Confusion Matrix Visualization



The confusion matrix shows prediction patterns across all 15 mammal classes. Diagonal values represent correct predictions, while off-diagonal values indicate misclassifications.

## 9.3 Per-Class Performance Analysis

| Class            | Precision   | Recall      | F1-Score    | Support     |
|------------------|-------------|-------------|-------------|-------------|
| bear             | 0.50        | 0.51        | 0.51        | 100         |
| fox              | 0.59        | 0.53        | 0.56        | 100         |
| hamster          | 0.59        | 0.60        | 0.59        | 100         |
| leopard          | 0.52        | 0.56        | 0.54        | 100         |
| lion             | 0.54        | 0.67        | 0.60        | 100         |
| mouse            | 0.44        | 0.41        | 0.42        | 100         |
| opossum          | 0.87        | 0.84        | 0.85        | 100         |
| porcupine        | 0.89        | 0.90        | 0.90        | 100         |
| rabbit           | 0.43        | 0.36        | 0.39        | 100         |
| raccoon          | 0.61        | 0.56        | 0.58        | 100         |
| shrew            | 0.44        | 0.44        | 0.44        | 100         |
| skunk            | 0.81        | 0.79        | 0.80        | 100         |
| squirrel         | 0.37        | 0.41        | 0.39        | 100         |
| tiger            | 0.68        | 0.52        | 0.59        | 100         |
| wolf             | 0.53        | 0.66        | 0.59        | 100         |
| <b>Macro Avg</b> | <b>0.59</b> | <b>0.58</b> | <b>0.58</b> | <b>1500</b> |

## 9.4 Key Observations from Confusion Matrix

### Best Classified Classes:

- **Porcupine** (90% recall, 89% precision) - Distinctive spiny appearance
- **Opossum** (84% recall, 87% precision) - Unique body shape and facial features
- **Skunk** (79% recall, 81% precision) - Distinctive black and white coloring

### Most Confused Classes:

- **Rabbit** (36% recall) - Often confused with other small mammals
- **Mouse** (41% recall) - Very similar to shrew and hamster
- **Squirrel** (41% recall) - Confused with similar-sized rodents

### Analysis:

1. Animals with distinctive visual patterns (porcupine, skunk, opossum) achieve highest accuracy
2. Small mammals (mouse, shrew, rabbit, squirrel) show significant confusion due to similar size and shape
3. Large carnivores (leopard, lion, tiger) show moderate confusion due to similar feline body structure

4. The model successfully learns species-specific features but struggles with subtle inter-class differences within similar animal groups

## 10. Conclusion

---

This project successfully demonstrated the systematic development and optimization of CNN models for fine-grained image classification. Starting from a baseline model achieving 49.07% accuracy, we progressively improved performance through:

1. **Regularization techniques** (Dropout + Batch Normalization) → 53.60%
2. **Architecture tuning** (increased depth and filters) → 56.80%
3. **Adding depth** (three convolutional layers) → **58.40%**

The best model achieved a **9.33% improvement** over the baseline, demonstrating that careful architecture design and proper regularization are essential for training effective CNNs on relatively small datasets.

The confusion matrix analysis revealed that the model performs best on visually distinctive animals (porcupine: 90%, opossum: 84%, skunk: 79%) while struggling with similar-looking species within the same family groups.

## Appendix: Code Repository Structure

---

```
HKUST_Python/
├── Python_project.ipynb      # Main Jupyter notebook with all code
├── report.md                 # This report
└── data/                      # CIFAR-100 dataset
```